



INSTITUTE OF TECHNOLOGY OF CAMBODIA

Department of Applied Mathematics and Statistics

Probabilistic Graphical Models

I4-AMS-DAS-A (2025–2026)

Building a Robust Part-of-Speech Tagger using Hidden Markov Model

Team 3

Name of Student	Student ID	Score
1. HEAM VISAL	e20221349	...

Lecturer: Mr. PHOK Ponna
Mr. PHOK Ponna

1 Abstract

This report presents a Part-of-Speech (POS) tagger based on a Hidden Markov Model (HMM). The system was implemented from scratch and evaluated on the Brown Corpus from NLTK using the universal POS tagset. The task is to assign the most likely POS tag to each word in a sentence. The dataset contains 57,340 sentences and was divided into 80% training data and 20% testing data. The HMM parameters were estimated using Maximum Likelihood Estimation from frequency counts. To handle unseen words, Laplace smoothing with $\alpha = 1.0$ and a special `<OOV>` token were used. The Viterbi algorithm was implemented in log-space to avoid numerical underflow. On the test set, the model tagged 217,777 out of 232,177 words correctly, giving an accuracy of 93.8%. The results show that a simple statistical sequence model can achieve strong performance for POS tagging, although errors remain for ambiguous and unseen words.

2 Introduction

Part-of-Speech tagging is a basic task in Natural Language Processing (NLP). It assigns a grammatical category, such as noun, verb, adjective, or adposition, to each word in a sentence. POS tagging is useful because it provides syntactic information for many later NLP tasks, including parsing, information extraction, machine translation, and text analysis.

This mini project builds a POS tagger using a Hidden Markov Model. An HMM is a probabilistic model for sequence data. In this task, the observed sequence is the sentence, and the hidden sequence is the POS tag sequence. The model assumes that each tag depends mainly on the previous tag, and that each word is generated from its tag.

The project uses the Brown Corpus from NLTK with the universal POS tagset. The universal tagset reduces the original detailed tag labels into 12 general POS tags. This makes the task simpler and allows the model to focus on common grammatical categories. The main objective is to train the HMM from the training sentences and use it to predict POS tags for unseen test sentences.

3 Methodology

3.1 Dataset and Tagset

The dataset used in this project is the Brown Corpus provided by NLTK. The corpus was loaded with the universal POS tagset. In total, the dataset contains 57,340 tagged sentences. The data was split into 45,872 training sentences and 11,468 testing sentences, following an 80%-20% split. The vocabulary size in the training data is 45,153, and the number of POS tags is 12.




Figure 1: Dataset split between training and testing sentences.

3.2 Hidden Markov Model

The POS tagging problem is modeled as a sequence prediction problem. Given a sentence $w_1, w_2, \dots, w_n$, the goal is to find the most likely tag sequence $t_1, t_2, \dots, t_n$. The HMM represents the joint probability of the tag and word sequence as:

$$P(t_1, \dots, t_n, w_1, \dots, w_n) = \pi(t_1) B_{t_1}(w_1) \prod_{i=2}^n A_{t_{i-1}, t_i} B_{t_i}(w_i) \quad (1)$$

In this formula, π is the initial probability, A is the transition probability matrix, and B is the emission probability matrix. These parameters are estimated from the training data.

3.3 Parameter Estimation

The model parameters were estimated using Maximum Likelihood Estimation based on frequency counts. The initial probability estimates how likely a tag is to start a sentence:

$$\pi(t) = \frac{\text{Count}(t \text{ starts a sentence})}{\text{Total sentences}} \quad (2)$$

The transition probability estimates how likely tag j is to follow tag i :

$$A(i, j) = \frac{\text{Count}(\text{tag}_i \text{ followed by tag}_j)}{\text{Count}(\text{tag}_i)} \quad (3)$$

The emission probability estimates how likely word w is generated by tag t . Laplace smoothing was used to reduce zero-probability problems:

$$B(t, w) = \frac{\text{Count}(w \text{ emitted by } t) + \alpha}{\text{Count}(t) + \alpha|V|} \quad (4)$$

In this project, $\alpha = 1.0$. A special $\langle \text{OOV} \rangle$ token was used for unseen words. During testing, words not found in the training vocabulary were treated as OOV words.

3.4 Decoding with Viterbi Algorithm

After estimating the HMM parameters, the Viterbi algorithm was used to find the best tag sequence for each test sentence. The algorithm uses dynamic programming to avoid checking all possible tag sequences directly. Since multiplying many small probabilities can cause computational underflow, the implementation uses log-space. The recurrence is:

$$\log v_t(j) = \max_i [\log v_{t-1}(i) + \log A(i, j)] + \log B(j, w_t) \quad (5)$$

The model stores the best previous tag for each current tag and then performs back-tracking to recover the predicted tag sequence.

4 Results and Evaluation

4.1 Evaluation Metric

The model was evaluated using word-level accuracy. Accuracy measures the proportion of test words whose predicted POS tag is equal to the true POS tag:

$$\text{Accuracy} = \frac{\text{Correctly tagged words}}{\text{Total test words}} \quad (6)$$

4.2 Final Evaluation Summary

Table 1 shows the final evaluation summary. The model achieved 93.8% accuracy on 232,177 test words. Out of these words, 217,777 were tagged correctly and 14,400 were tagged incorrectly.

Table 1: Final Evaluation Summary

Metric	Value
Training sentences	45,872
Testing sentences	11,468
Vocabulary size	45,153
POS tags	12
Total test words	232,177
Correct words	217,777
Wrong words	14,400
Accuracy	93.8%
OOV words	5,050
OOV rate	2.18%




Figure 2: Word-level POS tagging result on the test set.

4.3 OOV Analysis

The test set contains 5,050 OOV words, giving an OOV rate of 2.18%. OOV words are words that do not appear in the training vocabulary. The model handles these words using Laplace smoothing and the $\langle \text{OOV} \rangle$ token. This allows the model to assign non-zero emission probabilities to unseen words. However, OOV words are still difficult because the model has no direct training evidence for their most likely tags.




Figure 3: OOV word analysis on the test set.




Figure 4: Top 10 POS tagging mistake types.

5 Error Analysis

The model made 14,400 word-level mistakes on the test set. Many mistakes happen because some English words are ambiguous. A word can have different POS tags depending on how it is used in a sentence. For example, “average” can be a noun, adjective, or verb depending on context. Similarly, “offer” can be a noun or a verb.

Table 2 shows selected failed examples from the test sentences.

Table 2: Error Analysis Examples

Sentence	Word	True Tag	Predicted Tag
With tips, the girls average between \$150 and \$200 a week, depending on basic salary.	average	VERB	NOUN
	\$200 depending	NOUN ADP	ADP VERB
It was a very tempting offer.	very	ADV	ADJ
	tempting	VERB	NOUN
	offer	NOUN	VERB
He saw the surprise in her face, and laughed as though it were the funniest expression he had ever seen.	as	ADP	ADV

The HMM uses the first-order Markov assumption. This means that the current tag mainly depends on the previous tag, not on the full meaning of the sentence. Because of this, the model may fail when the correct tag requires longer context or semantic understanding. For example, deciding whether “offer” is a noun or a verb may require understanding its role in the sentence, not only the previous tag.

Rare words and unseen words are also challenging. Laplace smoothing helps by pre-

venting zero probabilities, but it cannot fully replace real evidence from training data. Therefore, OOV words can still receive incorrect tags, especially when they appear in unusual contexts.

6 Conclusion

This project implemented a POS tagger using a Hidden Markov Model trained on the Brown Corpus with the universal POS tagset. The model estimated initial, transition, and emission probabilities from frequency counts using Maximum Likelihood Estimation. Laplace smoothing with $\alpha = 1.0$ and a `<OOV>` token were used to handle unseen words. The Viterbi algorithm was implemented in log-space to decode the most likely POS tag sequence while avoiding computational underflow.

The final result shows that the HMM POS tagger achieved 93.8% accuracy on the test set. This is a strong result for a simple statistical model built from scratch. The main weaknesses are related to ambiguous words, limited context from the first-order Markov assumption, and rare or unseen words. Future improvements could include using richer word features, suffix information, higher-order tag dependencies, or modern neural sequence models.

References

- [1] Steven Bird, Ewan Klein, and Edward Loper. *Natural Language Processing with Python: Analyzing Text with the Natural Language Toolkit*. O'Reilly Media, 2009. Official online version available at <https://www.nltk.org/book/>.
- [2] NLTK Project. *NLTK Corpus Reader Documentation*. Official documentation showing access to the Brown Corpus and the universal POS tagset. Available at <https://www.nltk.org/howto/corpus.html>.
- [3] Andrew J. Viterbi. Error bounds for convolutional codes and an asymptotically optimum decoding algorithm. *IEEE Transactions on Information Theory*, 13(2):260–269, 1967. DOI: <https://doi.org/10.1109/TIT.1967.1054010>.
- [4] Lawrence R. Rabiner. A tutorial on hidden Markov models and selected applications in speech recognition. *Proceedings of the IEEE*, 77(2):257–286, 1989. DOI: <https://doi.org/10.1109/5.18626>.