



INSTITUTE OF TECHNOLOGY OF CAMBODIA

Department of Applied Mathematics and Statistics

Probabilistic Graphical Models

I4-AMS-DAS-A (2025–2026)

Building a Robust Part-of-Speech Tagger using Hidden Markov Model

Team 3

Name of Student	Student ID	Score
1. HEAM VISAL	e20221349	...
2. MON Sreylin	e20221701	...
3. CHOUB Botumraksa	e20221709	...
4. CHIV Mengchou	e20221028	...

Lecturer: Mr. PHOK Ponna

Contents

1	Abstract	1
2	Introduction	1
3	Methodology	1
3.1	Dataset and Tagset	1
3.2	Hidden Markov Model	1
3.3	Parameter Estimation	2
3.4	Decoding with Viterbi Algorithm	2
4	Results and Evaluation	3
4.1	Evaluation Metric	3
4.2	Final Evaluation Summary	3
4.3	OOV Analysis	3
5	Error Analysis	4
6	Conclusion	4

1 ABSTRACT

This report presents a Part-of-Speech (POS) tagger based on a Hidden Markov Model (HMM). The system was implemented from scratch and evaluated on the Brown Corpus from NLTK using the universal POS tagset. The task is to assign the most likely POS tag to each word in a sentence. The dataset contains 57,340 sentences and was divided into 80% training data and 20% testing data. The HMM parameters were estimated using Maximum Likelihood Estimation from frequency counts. To handle unseen words, Laplace smoothing with $\alpha = 1.0$ and a special $\langle 00V \rangle$ token were used. The Viterbi algorithm was implemented in log-space to avoid numerical underflow. On the test set, the model tagged 217,777 out of 232,177 words correctly, giving an accuracy of 93.8%. The results show that a compact statistical sequence model can recover most tag assignments from local sequential regularities. The remaining errors are concentrated around lexical ambiguity, sparse training evidence, and unseen words, which are limitations expected from the model assumptions and the chosen emission representation.

2 INTRODUCTION

Part-of-Speech tagging is a basic task in Natural Language Processing (NLP). It assigns a grammatical category, such as noun, verb, adjective, or adposition, to each word in a sentence. POS tagging is useful because it provides syntactic information for many later NLP tasks, including parsing, information extraction, machine translation, and text analysis.

This mini project builds a POS tagger using a Hidden Markov Model. An HMM is a probabilistic model for sequence data. In this task, the observed sequence is the sentence, and the hidden sequence is the POS tag sequence. The model assumes that each tag depends mainly on the previous tag, and that each word is generated from its tag.

The project uses the Brown Corpus from NLTK with the universal POS tagset. The universal tagset reduces the original detailed tag labels into 12 general POS tags. This makes the task simpler and allows the model to focus on common grammatical categories. The main objective is to train the HMM from the training sentences and use it to predict POS tags for unseen test sentences.

The work follows a complete supervised sequence-labeling pipeline. Tagged sentences provide the counts needed to estimate the model, a held-out portion of the corpus measures generalization, and Viterbi decoding converts the learned probabilities into one globally consistent tag sequence per sentence. This is important because selecting the most frequent tag for each word

independently would ignore the way grammatical categories tend to follow one another.

The HMM also provides an interpretable baseline. Its initial, transition, and emission components have direct probabilistic meanings, so a prediction can be related to sentence-start patterns, neighboring tags, and word-tag evidence. At the same time, the model deliberately uses limited context. Evaluating it therefore shows both how much can be learned from first-order dependencies and where richer lexical or contextual information would be needed.

3 METHODOLOGY

3.1 Dataset and Tagset

The dataset used in this project is the Brown Corpus provided by NLTK. The corpus was loaded with the universal POS tagset. In total, the dataset contains 57,340 tagged sentences. The data was split into 45,872 training sentences and 11,468 testing sentences, following an 80%–20% split. The vocabulary size in the training data is 45,153, and the number of POS tags is 12.

The sentence list was shuffled with random seed 42 before the split, making the experiment reproducible while keeping complete sentences together. No test sentence contributes counts to parameter estimation. This separation matters: the reported accuracy measures predictions on held-out sentences rather than the model's ability to reproduce the observations from which its probabilities were calculated. Words are normalized to lowercase for vocabulary lookup and emission counting, so capitalization variants share the same lexical evidence.




Figure 1: Dataset split between training and testing sentences.

3.2 Hidden Markov Model

The POS tagging problem is modeled as a sequence prediction problem. Given a sentence $w_1, w_2, \dots, w_n$, the goal is to find the most likely tag sequence $t_1, t_2, \dots, t_n$. The HMM represents the joint probability of the tag and word sequence as:

$$P(t_1, \dots, t_n, w_1, \dots, w_n) = \pi(t_1) B_{t_1}(w_1) \prod_{i=2}^n A_{t_{i-1}, t_i} B_{t_i}(w_i) \quad (1)$$

In this formula, π is the initial probability, A is the transition probability matrix, and B is the emission probability matrix. These parameters are estimated from the training data.

The factorization makes two simplifying assumptions. First, the tag at position i is conditioned on the tag at position $i - 1$ rather than on the entire tag history. Second, once the current tag is known, the observed word is treated as independent of the other words and tags. Under these assumptions, $\pi(t_1)$ scores a possible sentence opening, each A_{t_{i-1}, t_i} scores a local tag change, and each $B_{t_i}(w_i)$ scores the compatibility of a word with its tag. Multiplying the factors assigns a score to a complete word-tag sequence.

Although the hidden tags are available during supervised training, they are not available when a new sentence is decoded. At prediction time the words are the observations and the POS labels are hidden states that must be inferred jointly. This separation between observed words and latent test-time states is the reason the tagging task fits the HMM formulation.

3.3 Parameter Estimation

The model parameters were estimated using Maximum Likelihood Estimation based on frequency counts. The initial probability estimates how likely a tag is to start a sentence:

$$\pi(t) = \frac{\text{Count}(t \text{ starts a sentence})}{\text{Total sentences}} \quad (2)$$

The transition probability estimates how likely tag j is to follow tag i :

$$A(i, j) = \frac{\text{Count}(\text{tag}_i \text{ followed by tag}_j)}{\text{Count}(\text{tag}_i)} \quad (3)$$

The emission probability estimates how likely word w is generated by tag t . Laplace smoothing was used to reduce zero-probability problems:

$$B(t, w) = \frac{\text{Count}(w \text{ emitted by } t) + \alpha}{\text{Count}(t) + \alpha(|V| + 1)} \quad (4)$$

Here, V is defined as the set of distinct normalized words observed in the training data; it excludes the special $\langle \text{OOV} \rangle$ token. The additional

1 in $|V| + 1$ accounts for that OOV outcome. In this project, $\alpha = 1.0$. During testing, words not found in V were mapped to the OOV token. Its emission probability for tag t is therefore

$$B(t, \text{OOV}) = \frac{\alpha}{\text{Count}(t) + \alpha(|V| + 1)}. \quad (5)$$

The three estimates summarize different count tables. Initial counts are collected only from the first token of each training sentence. Transition counts are collected from adjacent tag pairs within a sentence, so sentence boundaries do not create artificial transitions. Emission counts are collected from normalized word-tag pairs. Dividing each count by its appropriate total turns the table into a conditional probability distribution.

Laplace smoothing is applied to emissions over the $|V|$ observed training words plus the one additional OOV outcome. It gives every word-tag combination a non-zero probability, while combinations observed many times still receive greater probability mass. The initial and transition distributions are estimated without Laplace smoothing in this implementation. When decoding, a small numerical constant is added before taking logarithms so that a zero estimate does not cause an undefined computation. This numerical safeguard should not be confused with re-estimating those distributions with smoothing.

3.4 Decoding with Viterbi Algorithm

After estimating the HMM parameters, the Viterbi algorithm was used to find the best tag sequence for each test sentence. The algorithm uses dynamic programming to avoid checking all possible tag sequences directly. Since multiplying many small probabilities can cause computational underflow, the implementation uses log-space. For a candidate tag j , initialization combines the sentence-start and first-word probabilities:

$$\log v_1(j) = \log \pi(j) + \log B(j, w_1). \quad (6)$$

For positions $t = 2, \dots, n$, the recurrence keeps the best score ending in j :

$$\log v_t(j) = \max_i [\log v_{t-1}(i) + \log A(i, j)] + \log B(j, w_t). \quad (7)$$

At the same time, the backpointer records which previous tag attained that maximum:

$$\psi_t(j) = \arg \max_i [\log v_{t-1}(i) + \log A(i, j)]. \quad (8)$$

The best final tag is selected from the last Viterbi column, after which the stored backpointers reconstruct the path in reverse:

$$\begin{aligned}\hat{t}_n &= \arg \max_j \log v_n(j), \\ \hat{t}_{t-1} &= \psi_t(\hat{t}_t), \quad t = n, n-1, \dots, 2.\end{aligned}\tag{9}$$

The model stores the best previous tag for each current tag and then performs backtracking to recover the predicted tag sequence.

Initialization combines the initial probability of each possible first tag with the first word’s emission probability. For every later position, the recurrence examines all previous tags and retains only the highest-scoring path ending in each current tag. After the final word, the best final state is selected and the stored backpointers are followed in reverse. With n words and T tags, this requires $O(nT^2)$ score comparisons instead of enumerating T^n complete tag sequences.

Log-space changes products into sums and preserves the ordering of path probabilities because the logarithm is monotonic. It therefore produces the same maximum-probability path as direct multiplication, subject to numerical precision, but avoids the rapid underflow caused by multiplying many values smaller than one.

4 RESULTS AND EVALUATION

4.1 Evaluation Metric

The model was evaluated using word-level accuracy. Accuracy measures the proportion of test words whose predicted POS tag is equal to the true POS tag:

The evaluation iterates through every token in the 11,468 held-out sentences. Each prediction is compared at the same sentence position with its corpus tag; correct comparisons are summed and divided by the total number of test tokens. This micro-averaged measure gives every word equal weight, regardless of its sentence length or tag frequency. It is easy to interpret, but it does not show whether performance is distributed uniformly across tags, known words, and OOV words. The subsequent counts and error analysis provide that additional context.

$$\text{Accuracy} = \frac{\text{Correctly tagged words}}{\text{Total test words}}\tag{10}$$

4.2 Final Evaluation Summary

Table 1 shows the final evaluation summary. The model achieved 93.8% accuracy on 232,177 test words. Out of these words, 217,777 were tagged correctly and 14,400 were tagged incorrectly.

The displayed percentage is rounded: the underlying evaluation counts are the primary result and satisfy $217,777 + 14,400 = 232,177$. Thus, the summary table and bar chart report the same evaluation from complementary views. The training and testing sentence counts also agree with the dataset split in Figure 1; they are included here to make the experimental scale visible beside the final score.

Metric	Value
Training sentences	45,872
Testing sentences	11,468
Vocabulary size	45,153
POS tags	12
Total test words	232,177
Correct words	217,777
Wrong words	14,400
Accuracy	93.8%
OOV words	5,050
OOV rate	2.18%

Table 1: Final Evaluation Summary




Figure 2: Word-level POS tagging result on the test set.

4.3 OOV Analysis

The test set contains 5,050 OOV words, giving an OOV rate of 2.18%. OOV words are words that do not appear in the training vocabulary. The model handles these words using Laplace smoothing and the $\langle \text{OOV} \rangle$ token. This allows the model to assign non-zero emission probabilities to unseen words. However, OOV words are still difficult because the model has no direct training evidence for their most likely tags.

The 5,050 OOV tokens are counted at test time by checking normalized words against the 45,153-word training vocabulary. The remaining 227,127 test tokens are known to the model. Therefore, the OOV analysis measures vocabulary coverage rather than a separate test subset, and the OOV

and known counts together equal the total number of evaluated words.

For an unseen word, every tag receives its tag-specific smoothed OOV emission probability. The decoder must then rely more heavily on the initial probability and surrounding transition patterns because no observed word-tag count identifies the word directly. Smoothing keeps decoding possible, but it does not use spelling, prefixes, suffixes, or other morphological cues that might separate likely nouns, verbs, and adjectives.




Figure 3: OOV word analysis on the test set.

5 ERROR ANALYSIS

The model made 14,400 word-level mistakes on the test set. Many mistakes happen because some English words are ambiguous. A word can have different POS tags depending on how it is used in a sentence. For example, “average” can be a noun, adjective, or verb depending on context. Similarly, “offer” can be a noun or a verb.

Table 2 shows selected failed examples from the test sentences.

The aggregate mistake chart in Figure 4 shows that errors are not evenly distributed. The largest direction is VERB predicted as NOUN (1,337 cases), followed by NOUN predicted as DET (1,113) and NOUN predicted as ADJ (986). Direction matters: NOUN predicted as VERB has 637 cases and is a different error category from VERB predicted as NOUN. Together, the ten displayed categories contain 7,730 of the 14,400 wrong predictions, so a substantial remainder is spread across less frequent confusion types.

These failures directly reflect the first-order Markov assumption, under which each tag depends only on the preceding tag rather than the full sentence context. In “With tips, the girls average...”, “average” is ambiguous among NOUN, ADJ, and VERB. Local tag transitions favored NOUN, and that wrong path also affected nearby predictions such as “\$200” and “depending.” The words “tempting” and “offer” in “It was a very tempting offer.” are lexically ambiguous. Relying only on local transition and emission probabilities,

the HMM selected an incorrect sequence without understanding the sentence meaning. Finally, in “He saw the surprise... laughed as though...”, “as” can be ADP or ADV. Recognizing the construction “as though” requires longer context than this first-order HMM represents.

Rare words and unseen words are also challenging. Laplace smoothing helps by preventing zero probabilities, but it cannot fully replace real evidence from training data. Therefore, OOV words can still receive incorrect tags, especially when they appear in unusual contexts.

The examples also illustrate why a single wrong decision can affect nearby positions. Viterbi chooses a complete path, and a locally attractive transition can change the best continuation for later words. Consequently, the rows in Table 2 should be read as outcomes of sequence decoding, not as independent word classifications. The table is illustrative; the overall counts and top mistake types summarize the complete test set.

6 CONCLUSION

This project implemented a POS tagger using a Hidden Markov Model trained on the Brown Corpus with the universal POS tagset. The model estimated initial, transition, and emission probabilities from frequency counts using Maximum Likelihood Estimation. Laplace smoothing with $\alpha = 1.0$ and a <OOV> token were used to handle unseen words. The Viterbi algorithm was implemented in log-space to decode the most likely POS tag sequence while avoiding computational underflow.

The final result shows that the HMM POS tagger achieved 93.8% accuracy on the test set, corresponding to 217,777 correct predictions among 232,177 words. The result demonstrates that initial-tag tendencies, adjacent-tag transitions, and word emissions capture enough regularity to label most tokens correctly. Because the score is computed on the held-out 20% split, it evaluates generalization to sentences that did not provide the training counts.

The analysis also identifies clear boundaries of the approach. The test set contains 5,050 OOV tokens, for which the model has no word-specific training count, and the most frequent mistakes involve confusions among common grammatical categories. The first-order state assumption represents only the previous tag, while the emission model uses a normalized word identity without morphological or semantic features. These design choices explain why smoothing prevents failure on unseen input but cannot resolve every ambiguous use.

Future improvements should be evaluated with




Figure 4: Top 10 POS tagging mistake types.

Sentence	Word	True Tag	Predicted Tag
With tips, the girls average between \$150 and \$200 a week, depending on basic salary.	average	VERB	NOUN
	\$200 depending	NOUN ADP	ADP VERB
It was a very tempting offer.	very	ADV	ADJ
	tempting	VERB	NOUN
	offer	NOUN	VERB
He saw the surprise in her face, and laughed as though it were the funniest expression he had ever seen.	as	ADP	ADV

Table 2: Error Analysis Examples

the same fixed split so that their effect can be compared directly with this baseline. Candidate extensions include suffix or other word-form features for unknown words, higher-order tag dependencies for broader local context, and neural sequence models that learn contextual representations. Such extensions add complexity, so the current HMM remains useful as a transparent reference model whose probabilities, decoding steps, and error counts can be inspected directly.

REFERENCES

- [1] Steven Bird, Ewan Klein, and Edward Loper. *Natural Language Processing with Python: Analyzing Text with the Natural Language Toolkit*. O’Reilly Media, 2009. Official online version available at <https://www.nltk.org/book/>.
- [2] NLTK Project. *NLTK Corpus Reader Documentation*. Official documentation showing access to the Brown Corpus and the universal POS tagset. Available at <https://www.nltk.org/howto/corpus.html>.
- [3] Andrew J. Viterbi. Error bounds for convolutional codes and an asymptotically optimum decoding algorithm. *IEEE Transactions on Information Theory*, 13(2):260–269, 1967. DOI: <https://doi.org/10.1109/TIT.1967.1054010>.
- [4] Lawrence R. Rabiner. A tutorial on hidden Markov models and selected applications in speech recognition. *Proceedings of the IEEE*, 77(2):257–286, 1989. DOI: <https://doi.org/10.1109/5.18626>.